

```
!pip install pandas
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (2.2.2)
Requirement already satisfied: numpy>=1.23.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
```

```
import pandas as pd
import numpy as np
```

```
df = pd.read_csv("Churn_Modelling.csv")
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   RowNumber             10000 non-null  int64
 1   CustomerId            10000 non-null  int64
 2   Surname               10000 non-null  object
 3   CreditScore           10000 non-null  int64
 4   Geography             10000 non-null  object
 5   Gender               10000 non-null  object
 6   Age                  10000 non-null  int64
 7   Tenure               10000 non-null  int64
 8   Balance              10000 non-null  float64
 9   NumOfProducts        10000 non-null  int64
10   HasCrCard            10000 non-null  int64
11   IsActiveMember       10000 non-null  int64
12   EstimatedSalary      10000 non-null  float64
13   Exited               10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
df.shape
```

```
(10000, 14)
```

```
df.duplicated().sum()
```

```
np.int64(0)
```

```
df.isna().sum()
```

	0
RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

```
df['Exited'].value_counts()
```

	count
Exited	
0	7963
1	2037

```
df['Geography'].value_counts()
```

	count
Geography	
France	5014
Germany	2509
Spain	2477

```
df['Gender'].value_counts()
```

	count
Gender	
Male	5457
Female	4543

```
df.drop(columns=['RowNumber', 'CustomerId', 'Surname'],inplace=True)
```

```
df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
df = pd.get_dummies(df, columns=['Geography', 'Gender'], drop_first=True)
```

df

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	Geography_Germany	Geography
0	619	42	2	0.00	1	1	1	101348.88	1	False	
1	608	41	1	83807.86	1	0	1	112542.58	0	False	
2	502	42	8	159660.80	3	1	0	113931.57	1	False	
3	699	39	1	0.00	2	0	0	93826.63	0	False	
4	850	43	2	125510.82	1	1	1	79084.10	0	False	
...
9995	771	39	5	0.00	2	1	0	96270.64	0	False	
9996	516	35	10	57369.61	1	1	1	101699.77	0	False	
9997	709	36	7	0.00	1	0	1	42085.58	1	False	
9998	772	42	3	75075.31	2	1	0	92888.52	1	True	
9999	792	28	4	130142.79	1	1	0	38190.78	0	False	

10000 rows × 12 columns

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
!pip install scikit-learn
```

```
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.11/dist-packages (1.6.1)
Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (2.0.2)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.14.1)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (3.6.0)
```

```
from sklearn.model_selection import train_test_split
```

```
X = df.drop(columns=['Exited'])
y = df['Exited']
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2, random_state=42)
```

```
X_train.shape
```

```
(8000, 11)
```

```
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
!pip install tensorflow==2.18.0
```

```
Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (25.2.10)
Requirement already satisfied: gast!=0.5.0,!<0.5.1,!<0.5.2,>=0.2.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (0.2.0)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (0.2.0)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (18.1.1)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (3.4.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (24.2)
Requirement already satisfied: protobuf!=4.21.0,!<4.21.1,!<4.21.2,!<4.21.3,!<4.21.4,!<4.21.5,<6.0.0dev,>=3.20.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (2.32.3)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (2.32.3)
Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (75.2.0)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (1.17.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (3.0.1)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (4.13.1)
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (1.17.2)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (1.71.0)
Collecting tensorboard<2.19,>=2.18 (from tensorflow==2.18.0)
  Downloading tensorboard-2.18.0-py3-none-any.whl.metadata (1.6 kB)
Requirement already satisfied: keras>=3.5.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (3.8.0)
Requirement already satisfied: numpy<2.1.0,>=1.26.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (2.0.2)
Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (3.13.0)
Collecting ml-dtypes<0.5.0,>=0.4.0 (from tensorflow==2.18.0)
  Downloading ml_dtypes-0.4.1-cp311-cp311-manylinux_2_17_x86_64_manylinux2014_x86_64.whl.metadata (20 kB)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.18.0) (0.37.0)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from astunparse>=1.6.0->tensorflow==2.18.0) (0.44.0)
Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow==2.18.0) (13.9.4)
Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow==2.18.0) (0.0.8)
```

```
Requirement already satisfied: certifi<2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<2.32.0,>=2.31.0->tensorflow==2.18.0) (2025.1.1)
```

```
Requirement already satisfied: certifi<2021.4.17 in /usr/local/lib/python3.11/dist-packages (from requests>=2.22.0->tensorflow==2.10.0)
Requirement already satisfied: markdown<=2.6.8 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow==2.10.0)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow==2.10.0)
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow==2.10.0)
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.11/dist-packages (from werkzeug>=1.0.1->tensorboard<2.19,>=2.18->tensorflow==2.10.0)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow==2.10.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow==2.10.0)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0->tensorflow==2.10.0)
Downloading tensorflow-2.18.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (615.4 MB)
615.4/615.4 MB 2.1 MB/s eta 0:00:00
Downloading ml_dtypes-0.4.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (2.2 MB)
2.2/2.2 MB 43.1 MB/s eta 0:00:00
Downloading tensorboard-2.18.0-py3-none-any.whl (5.5 MB)
5.5/5.5 MB 71.3 MB/s eta 0:00:00
Installing collected packages: ml-dtypes, tensorboard, tensorflow
Attempting uninstall: ml-dtypes
  Found existing installation: ml_dtypes 0.5.1
  Uninstalling ml_dtypes-0.5.1:
    Successfully uninstalled ml_dtypes-0.5.1
Attempting uninstall: tensorboard
  Found existing installation: tensorboard 2.19.0
  Uninstalling tensorboard-2.19.0:
    Successfully uninstalled tensorboard-2.19.0
Attempting uninstall: tensorflow
  Found existing installation: tensorflow 2.19.0
  Uninstalling tensorflow-2.19.0:
    Successfully uninstalled tensorflow-2.19.0
Successfully installed ml-dtypes-0.4.1 tensorboard-2.18.0 tensorflow-2.18.0
```

```
!python --version
```

Python 3.11.11

```
from keras.models import Sequential
from keras.layers import Dense, Input

# Initialize the model
model = Sequential()

# Adding the input layer
# We use an Input layer to define the input shape (11 features)
# Input layer doesn't need to be explicitly added as the first Dense layer
# will automatically infer the input shape, but we can define it for clarity.
model.add(Input(shape=(11,))) # Here, 11 is the number of features (input shape)

# # Adding the hidden layer
# # 3 perceptrons (neurons) in this hidden layer, with a sigmoid activation function
# model.add(Dense(3, activation='sigmoid')) # The number of perceptrons is 3

# # Adding the output layer
# # Output layer with 1 perceptron and sigmoid activation (for binary classification)
# model.add(Dense(1, activation='sigmoid'))

#making changes to improve accuracy
model.add(Dense(11, activation='relu'))
#first hidden layer with 11 perceptrons, relu activation fn
model.add(Dense(11, activation='relu'))
#second hidden layer with 11 perceptrons, relu activation fn
model.add(Dense(1, activation='sigmoid'))
```

```
model.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 11)	132
dense_8 (Dense)	(None, 11)	132
dense_9 (Dense)	(None, 1)	12

Total params: 276 (1.08 KB)

Trainable params: 276 (1.08 KB)

```
model.compile(loss = 'binary_crossentropy', optimizer = 'Adam', metrics = ['accuracy'])
```

```
history = model.fit(X_train_scaled, y_train, epochs = 100, validation_split = 0.2)
```

```
#history will store the info during each stage of training
```



epoch 74/100
200/200 1s 4ms/step - accuracy: 0.8632 - loss: 0.3310 - val_accuracy: 0.8594 - val_loss: 0.3418
Epoch 75/100
200/200 1s 3ms/step - accuracy: 0.8727 - loss: 0.3222 - val_accuracy: 0.8587 - val_loss: 0.3413
Epoch 76/100
200/200 1s 3ms/step - accuracy: 0.8676 - loss: 0.3174 - val_accuracy: 0.8594 - val_loss: 0.3415
Epoch 77/100
200/200 1s 4ms/step - accuracy: 0.8665 - loss: 0.3238 - val_accuracy: 0.8619 - val_loss: 0.3417
Epoch 78/100
200/200 1s 3ms/step - accuracy: 0.8651 - loss: 0.3341 - val_accuracy: 0.8587 - val_loss: 0.3427
Epoch 79/100
200/200 1s 3ms/step - accuracy: 0.8707 - loss: 0.3207 - val_accuracy: 0.8612 - val_loss: 0.3415
Epoch 80/100
200/200 1s 3ms/step - accuracy: 0.8674 - loss: 0.3259 - val_accuracy: 0.8600 - val_loss: 0.3424
Epoch 81/100
200/200 1s 3ms/step - accuracy: 0.8631 - loss: 0.3244 - val_accuracy: 0.8587 - val_loss: 0.3424
Epoch 82/100
200/200 1s 4ms/step - accuracy: 0.8629 - loss: 0.3234 - val_accuracy: 0.8569 - val_loss: 0.3421
Epoch 83/100
200/200 2s 6ms/step - accuracy: 0.8667 - loss: 0.3252 - val_accuracy: 0.8581 - val_loss: 0.3417
Epoch 84/100
200/200 1s 4ms/step - accuracy: 0.8674 - loss: 0.3257 - val_accuracy: 0.8594 - val_loss: 0.3420
Epoch 85/100
200/200 1s 4ms/step - accuracy: 0.8672 - loss: 0.3190 - val_accuracy: 0.8600 - val_loss: 0.3415
Epoch 86/100
200/200 1s 3ms/step - accuracy: 0.8635 - loss: 0.3246 - val_accuracy: 0.8594 - val_loss: 0.3418
Epoch 87/100
200/200 1s 3ms/step - accuracy: 0.8639 - loss: 0.3218 - val_accuracy: 0.8575 - val_loss: 0.3443
Epoch 88/100
200/200 1s 3ms/step - accuracy: 0.8641 - loss: 0.3291 - val_accuracy: 0.8594 - val_loss: 0.3419
Epoch 89/100
200/200 2s 5ms/step - accuracy: 0.8672 - loss: 0.3286 - val_accuracy: 0.8575 - val_loss: 0.3429
Epoch 90/100
200/200 1s 3ms/step - accuracy: 0.8652 - loss: 0.3277 - val_accuracy: 0.8575 - val_loss: 0.3422
Epoch 91/100
200/200 1s 3ms/step - accuracy: 0.8711 - loss: 0.3170 - val_accuracy: 0.8581 - val_loss: 0.3440
Epoch 92/100
200/200 1s 4ms/step - accuracy: 0.8661 - loss: 0.3188 - val_accuracy: 0.8569 - val_loss: 0.3425
Epoch 93/100
200/200 2s 6ms/step - accuracy: 0.8706 - loss: 0.3184 - val_accuracy: 0.8581 - val_loss: 0.3429
Epoch 94/100
200/200 2s 7ms/step - accuracy: 0.8716 - loss: 0.3258 - val_accuracy: 0.8581 - val_loss: 0.3420
Epoch 95/100
200/200 2s 3ms/step - accuracy: 0.8718 - loss: 0.3135 - val_accuracy: 0.8556 - val_loss: 0.3428
Epoch 96/100
200/200 1s 3ms/step - accuracy: 0.8698 - loss: 0.3215 - val_accuracy: 0.8562 - val_loss: 0.3433
Epoch 97/100
200/200 1s 3ms/step - accuracy: 0.8653 - loss: 0.3218 - val_accuracy: 0.8550 - val_loss: 0.3420
Epoch 98/100
200/200 1s 3ms/step - accuracy: 0.8637 - loss: 0.3411 - val_accuracy: 0.8562 - val_loss: 0.3447
Epoch 99/100
200/200 1s 3ms/step - accuracy: 0.8662 - loss: 0.3288 - val_accuracy: 0.8581 - val_loss: 0.3432
Epoch 100/100
200/200 1s 3ms/step - accuracy: 0.8658 - loss: 0.3225 - val accuracy: 0.8600 - val loss: 0.3425


history.history



```
0.341519504/85555//,
0.3424181044101715,
0.34243810176849365,
0.3421393930912018,
0.341742604970932,
0.3420385718345642,
0.341548889875412,
0.34184274077415466,
0.344260573387146,
0.3419116139411926,
0.3429023027420044,
0.34224164485931396,
0.3440454602241516,
0.3425360918045044,
0.34287789463996887,
0.3420100808143616,
0.342782199382782,
0.3433056175708771,
0.34202370047569275,
0.3446817696094513,
0.34315043687820435,
0.34250980615615845}]}
```

```
import matplotlib.pyplot as plt
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
```

↗ [matplotlib.lines.Line2D at 0x791cfb377010]




model.layers[0].get_weights()

↗ [array([[0.10235199, -0.02729407, -0.29689068, 0.17194222, -0.18882293, 0.05254847, -0.02412846, 0.17464419, 0.03954822, 0.2192056, 0.2657583], [0.13614827, 0.06501141, 1.1472802, 0.46831146, 0.68798125, -0.9236476, 0.90113753, -0.22191288, -0.16493824, -0.25321862, 0.2642599], [-0.07170658, 0.18295506, -0.26123932, -0.12318897, 0.33628646, 0.04486928, -0.07410918, -0.01639474, 0.11084247, 0.13971083, 0.26739368], [-0.82583696, 0.7422567, -0.4203305, -0.09842348, 0.22384715, -0.15849143, 0.07685588, -0.30745456, -0.01586948, -0.7000992, -0.432479], [1.2519408, 1.0856526, -0.10618119, 0.27615508, 0.463923, 0.14059034, -1.0662065, -0.03953137, 0.02789063, -1.2422748, -0.1934734], [0.1558718, 0.04680955, -0.30236602, 0.0962111, 0.5484045, 0.06579883, -0.00803998, -0.03814459, -0.17619398, 0.14231792, -0.0965118], [-0.13955241, 0.16007546, 0.8212532, 0.51250297, 0.44402683, 0.16252552, -0.44072834, -0.8481122, -0.22634505, 0.17889743, 0.80304], [-0.10055321, 0.13099238, 0.27322102, 0.1343742, -0.14304921, 0.1666675, -0.16194543, -0.11800639, -0.2868092, 0.08978967, -0.19111054], [0.30725777, -0.40322182, -0.00328511, 0.65779805, -0.21778814, 0.10186622, 0.07185046, 0.3972196, -0.91652155, 0.5378111, -0.7817219], [0.16286866, -0.01699267, -0.15994553, 0.5024812, 0.35650083, 0.08777893, -0.1793953, -0.3110754, -0.00772641, 0.2640108, -0.69570595], [0.02315623, -0.06625043, 0.6594553, 0.300621, 0.04006233, 0.20845628, 0.02255104, -0.15952115, 0.56285936, -0.01202548, -0.37703884]], dtype=float32), array([-0.30076805, -0.03947368, -0.873603, 0.37505266, -0.49558318, 1.1343312, 0.15990146, 0.6439129, 0.49139103, -0.18615898, -0.169265], dtype=float32)]

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
```

[<matplotlib.lines.Line2D at 0x791cfb20b390>]



```
y_log = model.predict(X_test_scaled)
```

63/63 ————— 0s 4ms/step

```
#The above results are probabilities of 'Exited' = 1
#We need to decide a threshold using ROC or other methods, for now, we will assume a threshold of 0.5
y_pred = np.where(y_log>=0.5, 1,0)
y_pred
```

```
array([[0],
       [0],
       [0],
       ...,
       [1],
       [0],
       [0]])
```

```
from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_pred)
```

0.863